

Distributed Systems Study Group

First Meeting: Dec 3rd, 2019

Gerry Seidman
gerry@iamx.com

Distributed Systems Study Group

- Get a deeper understanding of Distributed Systems
 - Architectural Patterns
 - Today's "standard" technology
 - Technology 'Internals'
- Collaborate with others to get Software installed/running
- Design and Build a Distributed Application from Scratch
- Maintain a public/shared Journal of what we have learned

Who are we?

- Developer
- Architect
- Admin
- SRE
- Application Security
- Data Scientist

This is a Study Group

- Everyone has something to contribute
 - Prepare and Give presentations
 - Help others in the group with their work
 - Take/Post notes to be posted onto our Shared Journal (git repo)
- Format
 - Hands-on / Networking
 - Questions for the Group
 - Presentation
 - Hands-on
- Not everyone will go the same speed
 - Each time we meet there will be hands-on time to help people catch up
 - You should always be able to follow the presentations

Areas that We can Cover

- Containers / Kubernetes
 - Application packaging, deployment and management
- Kafka
 - Used for building real-time data/event pipelines and streaming apps.
- Prometheus
 - Time Series Database used for event monitoring and alerting
- Grafana
 - Front end Dashboards for Monitoring and Analytics
- CI/CD (Jenkins, other)
 - Continuous Integration / Continuous Delivery
- Distributed Tracing
 - Used for building real-time data/event pipelines and streaming apps.
- Hadoop/Spark
- ELK Stack –Elastic Logstash Kibana
- Agile
 - JIRA

And for all of them...

Security ... oh Security

Goals for Hands-On Tonight

- VMs (or external machines) set up with the appropriate network inter-connectivity
- Install Linux (Headless, ie no desktop interface, only ssh)
- Have docker properly installed
- Have a kubernetes cluster installed

It Takes a Village

- Sharing/memorializing our experience

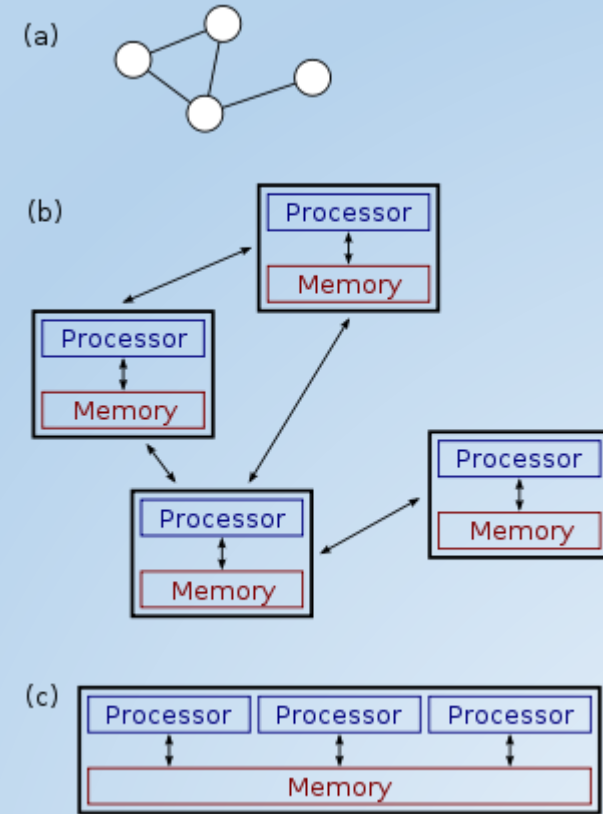
<https://github.com/GerrySeidman/Distributed-Systems-Study-Group>

Agenda (approximately)

- Overview of Distributed Systems
 - What we'd like to build
- Quick Overview
 - File Systems
 - Networking
- Understanding Containers
- Understanding Kubernetes

What is a Distributed System

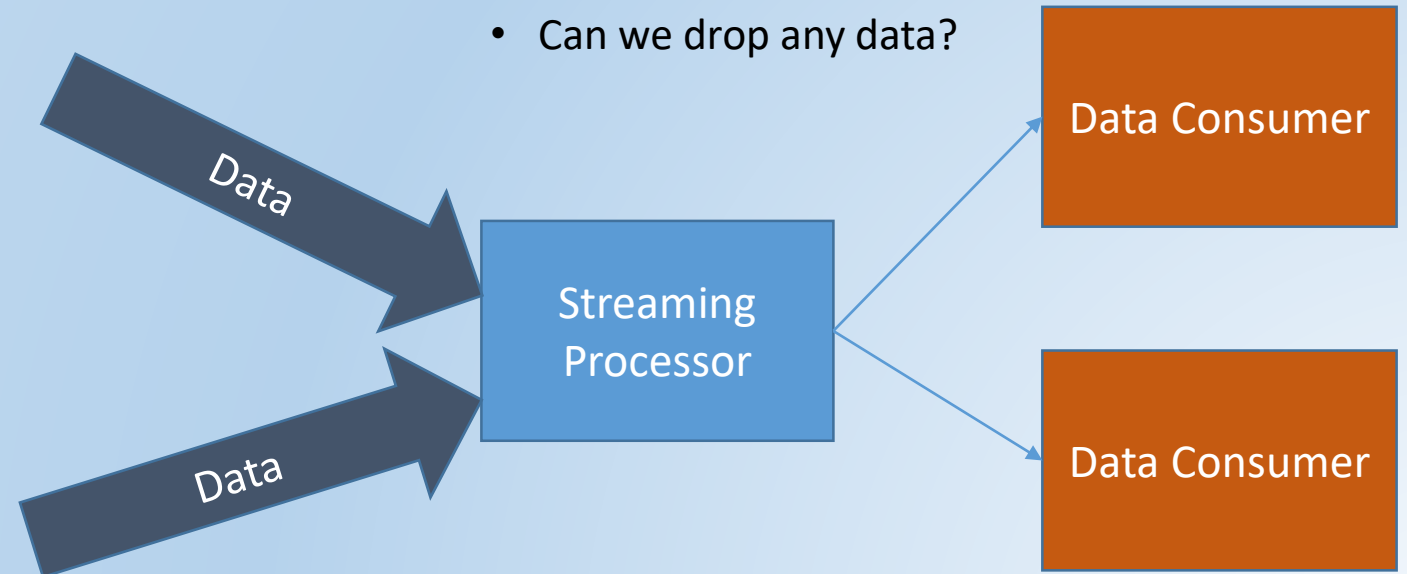
- Cluster Service
 - Database
 - Message Bus
 - HTTP Server
- Complex System
 - Many of the above
 - Working Together (loosely coupled)
 - Possibly spanning multiple data centers
 - Either Enterprise or Cloud
 - Buzz Words: Private Cloud, Multi-Cloud, Hybrid Cloud



For Example: How to handle Streaming data?

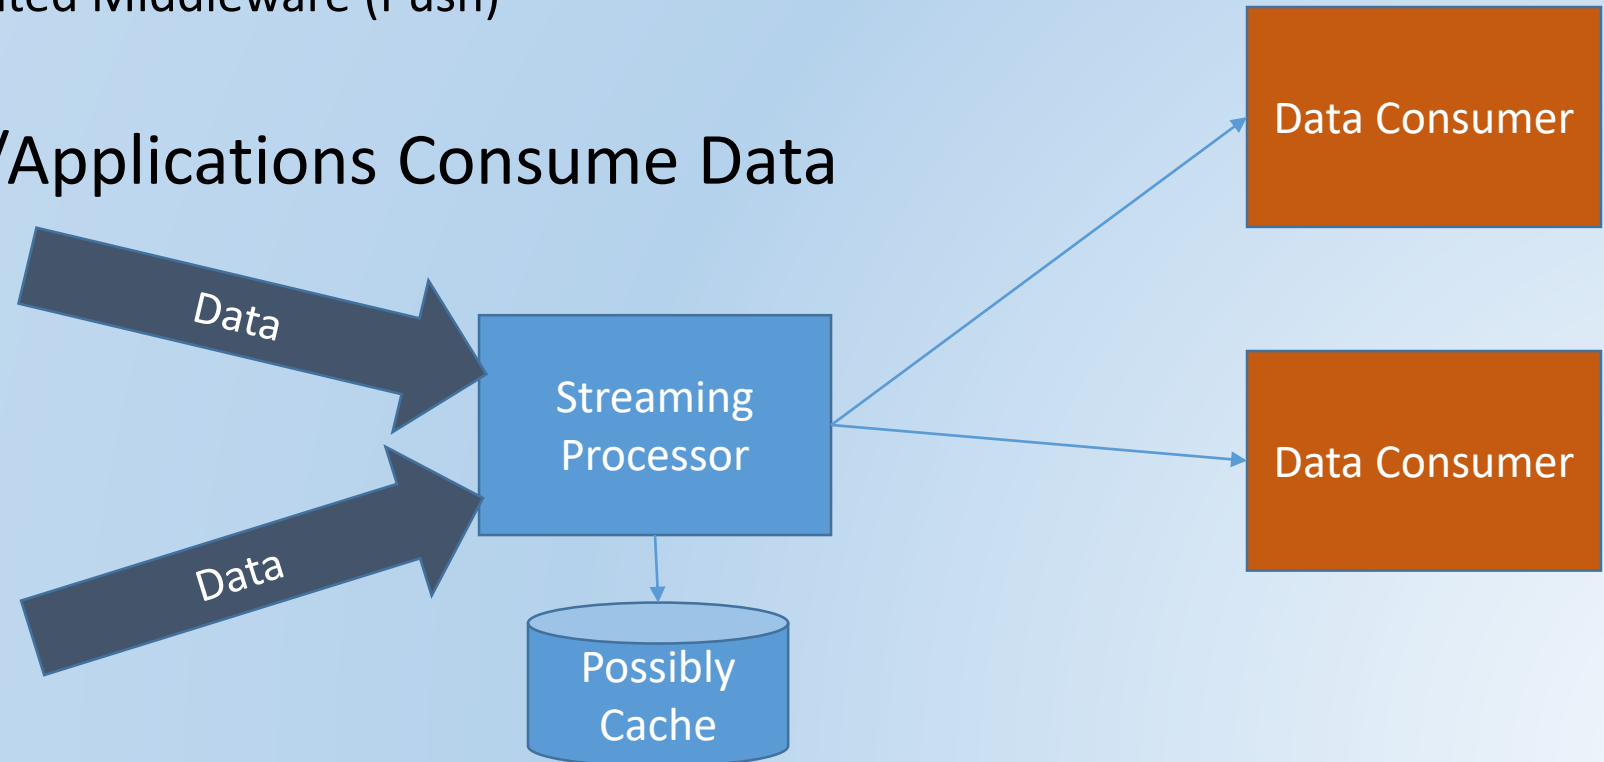
- Sources
 - Twitter Feeds
 - Stock Trades
 - Logs
 - SIEMS
- Possibly many
 - Homogeneous
 - Heterogenous
- Reliability
 - Availability
 - Scale on inbound data
 - Sustained or Bursty
 - Scale on # of Consumers
 - Can we drop any data?

- Consumers
 - Homogeneous
 - Heterogeneous
- Reliability
 - Scale on data that I want
 - Can we drop any data?

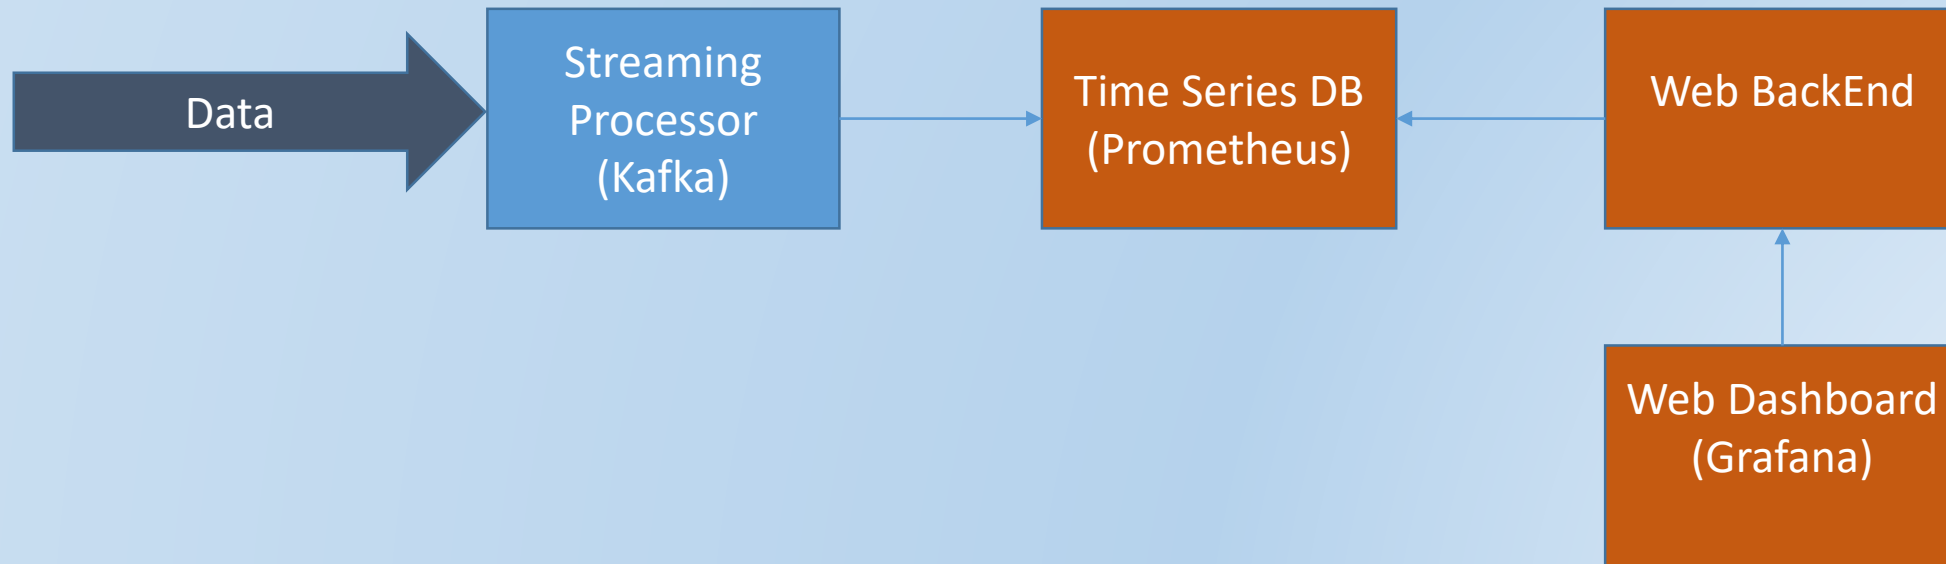


Streaming Data Processing

- Cluster of Streaming processors receive data
 - Persistence Model and Guarantees
 - Flavors
 - Message Oriented Middleware (Push)
 - Kafka (Pull)
- Other Processes/Applications Consume Data

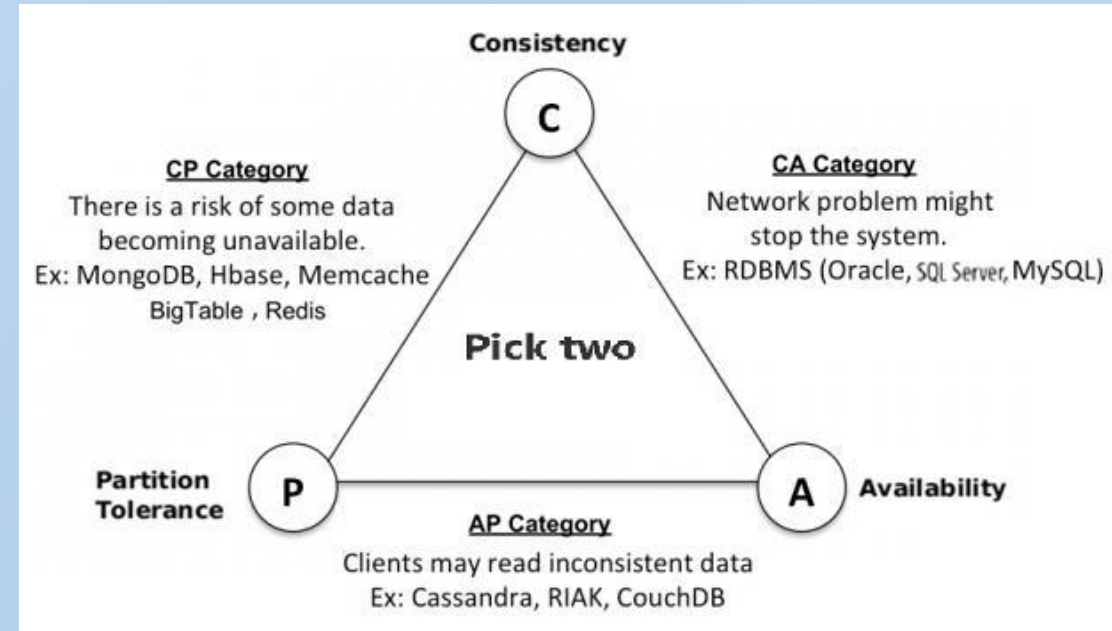


Complex Distributed System Example



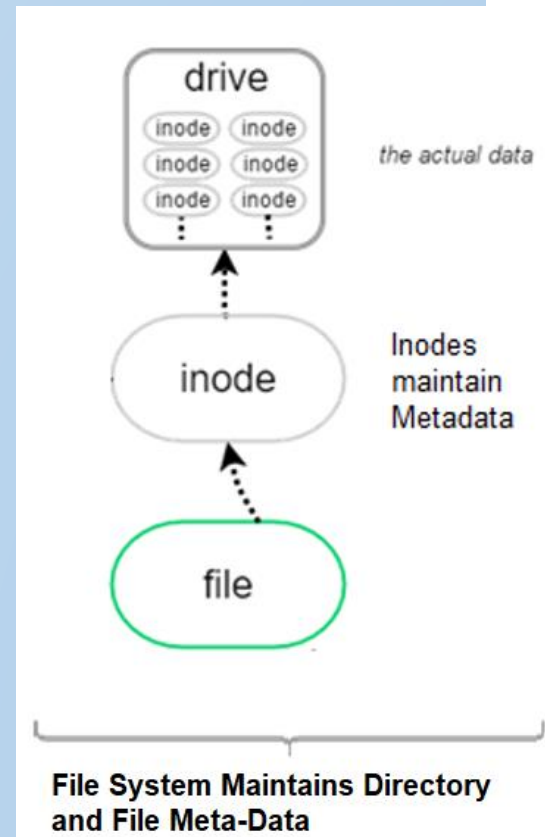
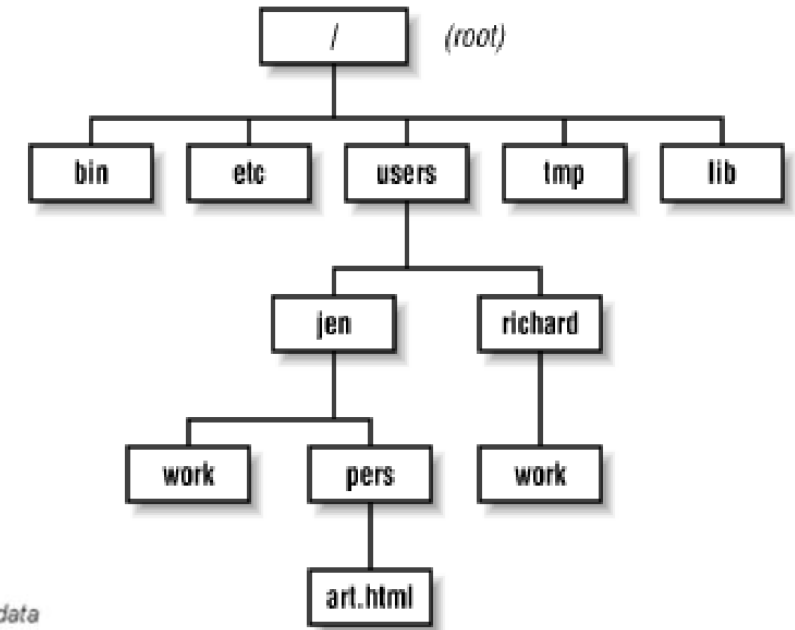
CAP Theorem - Tradeoffs

- Consistency:
 - Every read receives the most recent write or an error
- Availability:
 - Every request receives a (non-error) response, without the guarantee that it contains the most recent write
- Partition tolerance:
 - The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes



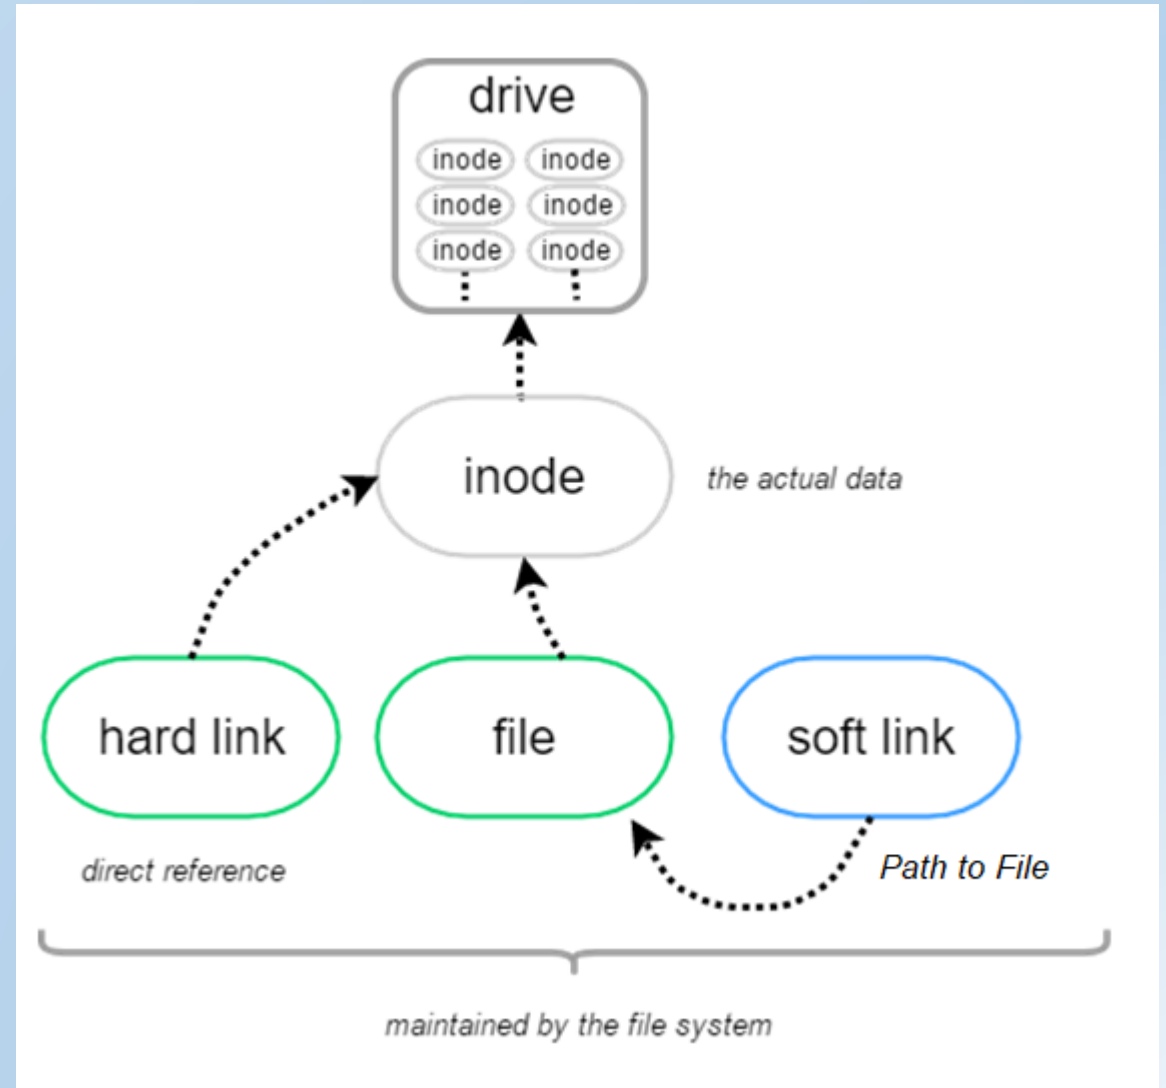
File Systems

- Volume (Logical/Physical Drives)
 - HD, SSD, CD, USB, Memory, Network
- Directory Structure
- File System
 - Linux EXT4, XFS, ZFS
 - Windows: FAT, NTFS
- Inodes
- Metadata + Extended Attributes



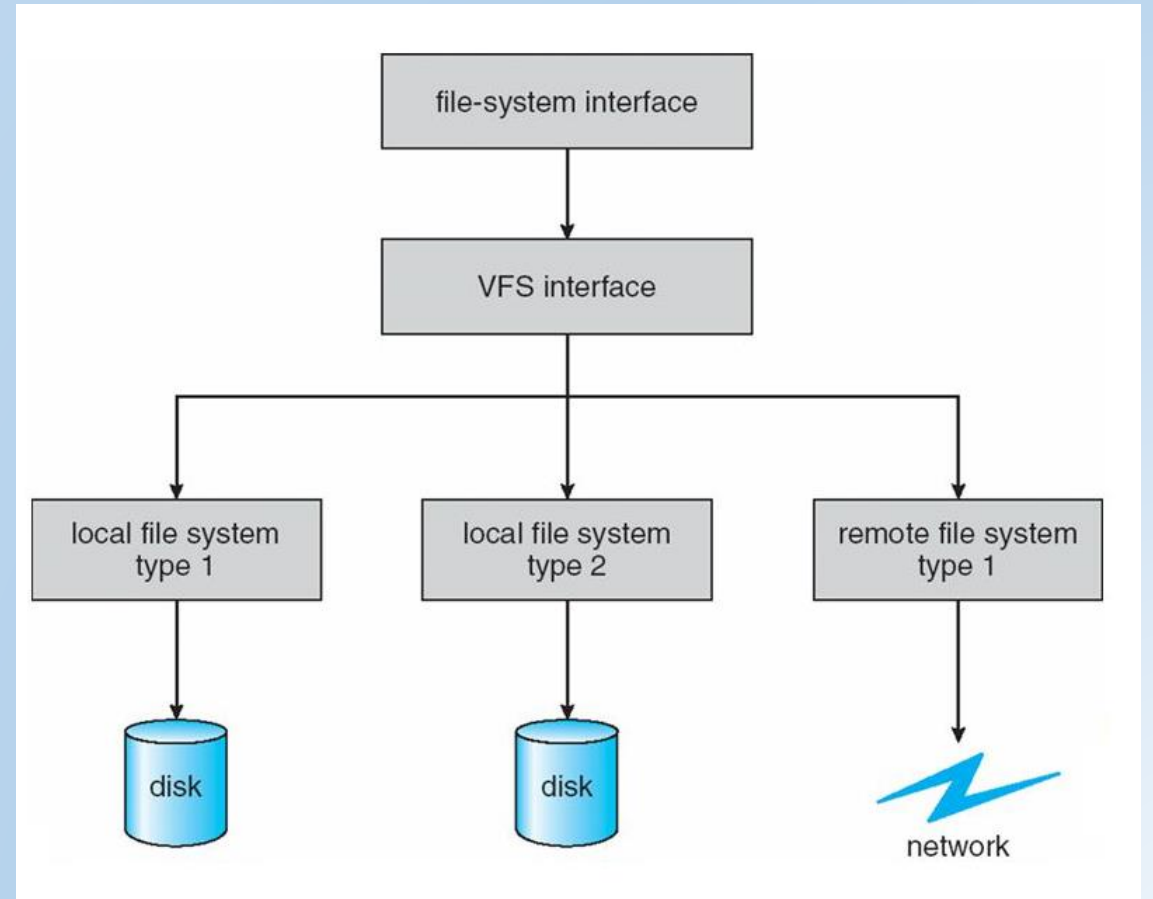
Links

- Hard
- Soft/Symbolic
- Works for Files or Directories



Virtual File System (VFS)

- Mounts
 - 'Mount' a FileSystem to a directory path in the VFS



Overlay File System

- We'll get Back to this with Containers....
- ~~Copy on Write~~
- ~~<https://docs.docker.com/storage/>~~
- ~~<https://docs.docker.com/storage/storagedriver/overlayfs-driver/>~~

I/O Basics

- Stream I/O
 - File
 - Pipe
 - Sockets
- Random Access I/O
 - Typically things like databases
 - But Databases may use raw Volumes directly

IP Addresses and DNS

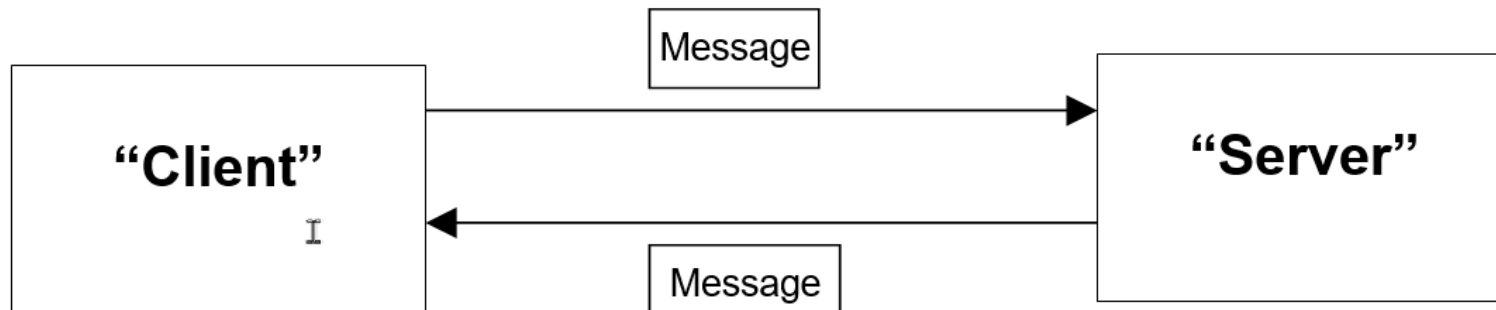
- IP Address
- DNS
- Subnet
 - Classless Inter-Domain Routing (CIDR)
 - Superseded Class only Routing ~1981 A/B/C/D
 - Subnet Designation Notation
 - 192.168.16.0/24
 - Alternatively
 - Network Mask: 255.255.255.0

Socket Basics

- An “server” application
 - Can register a ‘Service’
 - At a Port (1-65535)
 - On one of its Network Interface
 - Which can then ‘accept’ connections
 - Upon accept it has a Socket

- A “client” application
 - Can connect to a ‘Service’
 - On a Network Interface
 - At a Port (1-65535)
 - Upon Connect it has a Socket


- Each Side has a Socket
 - From then on it is all about ‘Protocol’



IP and TCP/IP

- IP Packets
 - Source Address/Port → Destination Address/Port + Payload
 - TCP/IP and UDP (others like ICMP for Ping)
- TCP Sockets
 - Frame Buffer
 - Writes can block

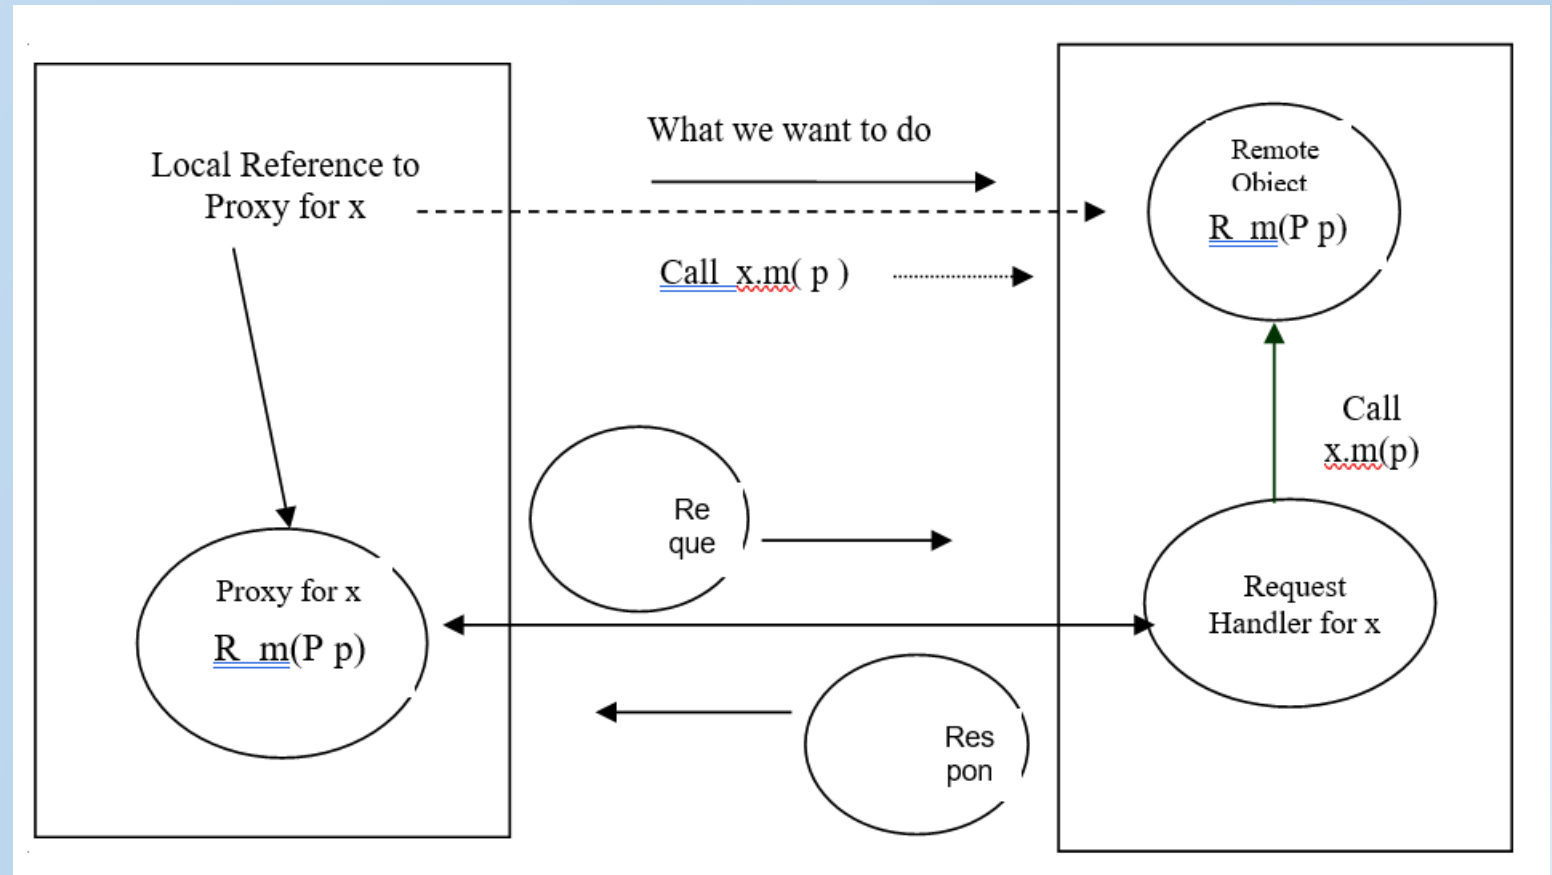
Calling a Function/Method on an Object



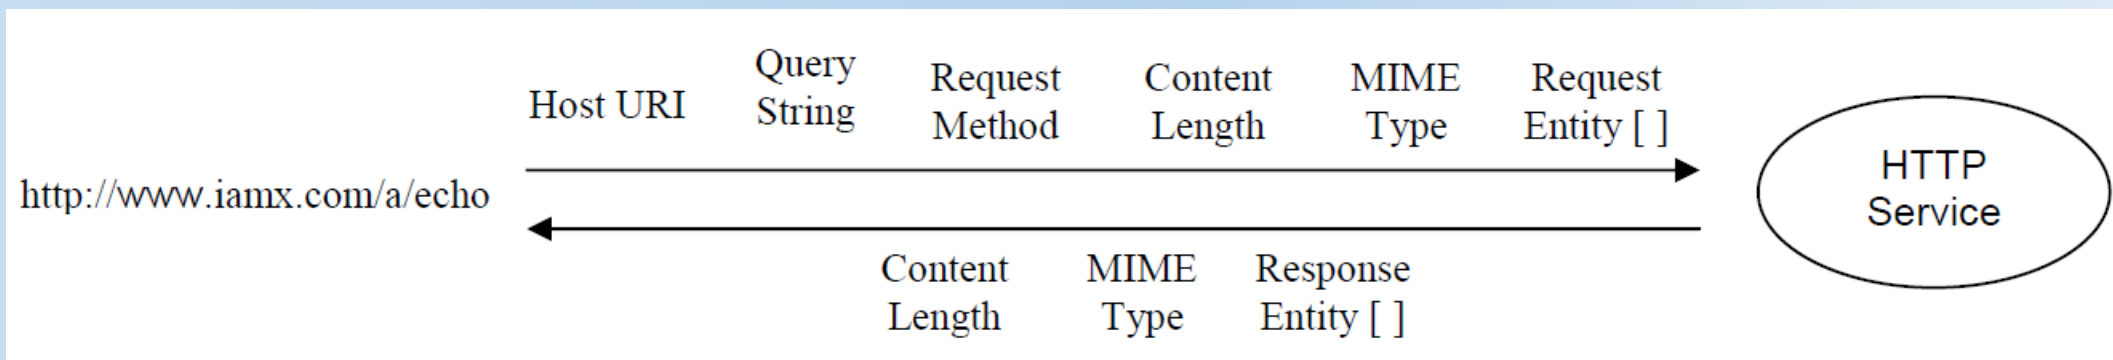
Remote Procedure Calls

- RPC Protocols

- REST
- gRPC
- CORBA



- Example REST over HTTP



Understanding HTTP Sessions

- `curl -v --cookie "Elf=Keebler" "http://www.reunion.allaire.com/somepath?dog=Lassie&cow=Elsie"`

* Connected to www.reunion.allaire.com (198.185.159.144) port 80 (#0)

> GET /somepath?dog=Lassie&cow=Elsie HTTP/1.1

> Host: www.reunion.allaire.com

> User-Agent: curl/7.55.1

> Accept: */*

> **Cookie: Elf=Keebler**

>

< HTTP/1.1 404 Not Found

< date: Tue, 03 Dec 2019 20:58:40 GMT

< expires: Thu, 01 Jan 1970 00:00:00 GMT

< content-type: text/html; charset=utf-8

< Age: 4

< **Set-Cookie: crumb=BXmfdu9zkjHqNmIzNGFhYWQ1NGIzZmU5YWFiNGNhNmU0NjVmYmM0;Path=/**

< Content-Length: 2052

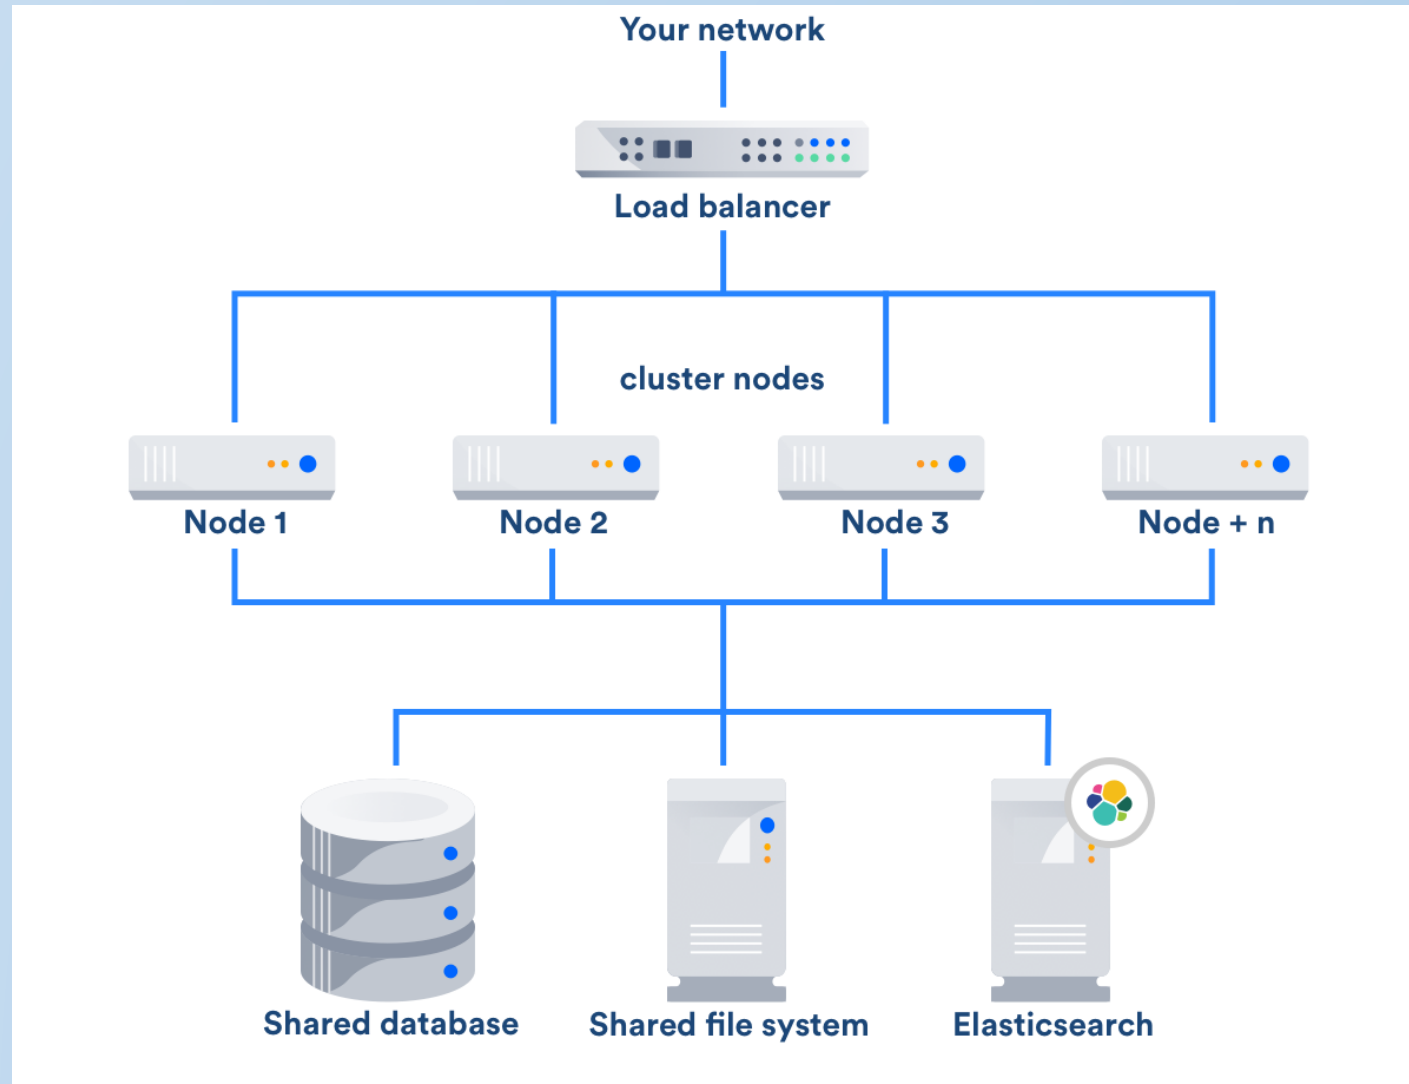
< x-contextid: 9toQ3jNh/GuREc2eU

< server: Squarespace

NAT, Firewalls, IP Tables

- Network Address Translation (NAT)
 - Uses a single external IP Address
 - Internally provides a Subnet (ie many IP addresses from one)
- IP Tables
 - Redirect or block packets
- Firewall Rules

Load Balancers and Content Switches

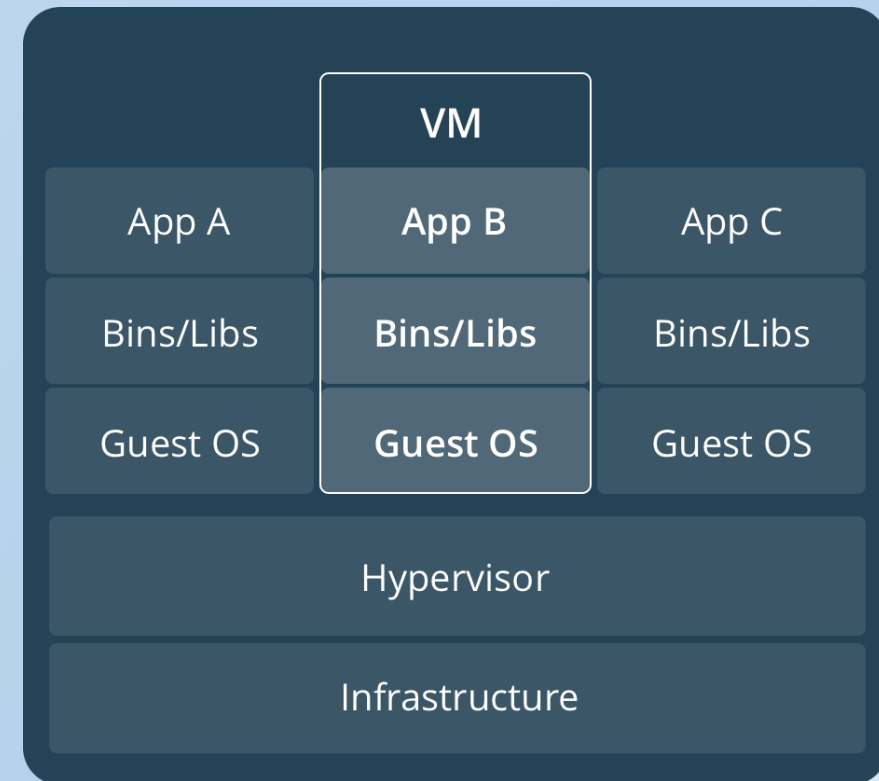


Software Packaging and Deployment

- Developer writes some code, tests it locally and we're ready to deploy
 - But where?
 - How to deliver it?
- Zip it up and install it on some dedicated machine
 - Oh yeah.. There's also configuration files
- You need a machine to handle peak load of that machine
 - How much Disk? RAM? Network? GPU? Fast Processor?
 - Hard Allocation of Resources



Virtual Machine

- Unit of deployment a VM Image
 - Just a big TAR file
 - When running it is a VM
- Semi-Soft Allocation of Resources
 - Disk, Memory, Network
- Slow to load
- Strong Security Boundary
- Hard to maintain
 - Same amount of work as a dedicated Machine
 - Security patches



Container

- Unit of deployment a Container Image
 - Just a big TAR file
 - When running it is a Container
- Soft Allocation of Resources
 - Disk, Memory
 - Some shared across Containers
- Fast to load
 - Run as a Host Process
- Pretty Strong Security Boundary
 - Linux Kernel
- Easier to maintain
 - Host OS

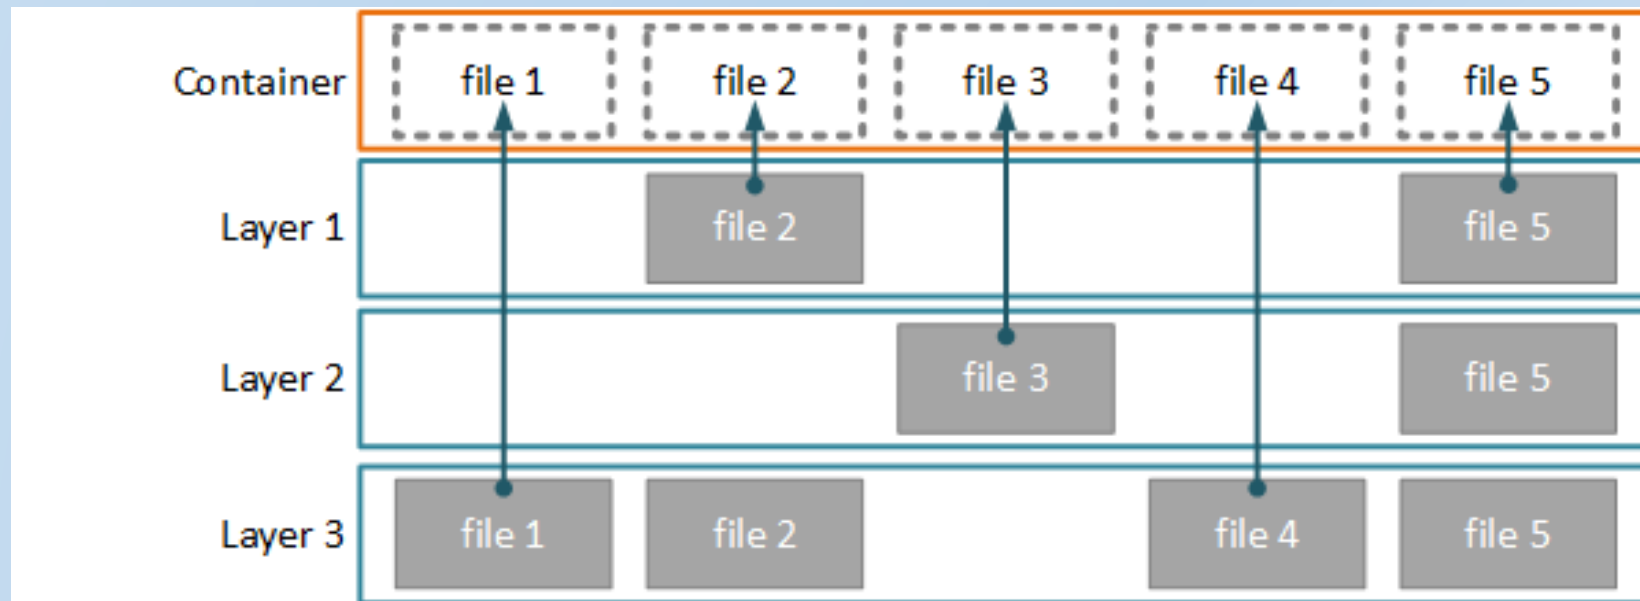


What is Docker

- Docker File Format for Container Images
- Docker Container Image builder
 - Must have Root access to build
- Docker Runtime 'server'
 - dockerd
 - REST or CLI Client
 - Must have Root access to launch containers
- Docker Registry
- Other Container implementations
 - Singularity
 - Podman/Buildah

Overlay File System – Container Image Layers

- Copy on Write File System



Container Resource Management

- Container Isolation
 - Control Groups (cgroups)
 - Namespaces

Namespaces

- The pid namespace: Process isolation (PID: Process ID).
- The net namespace: Managing network interfaces (NET: Networking).
- The ipc namespace: Managing access to IPC resources (IPC: InterProcess Communication).
- The mnt namespace: Managing filesystem mount points (MNT: Mount).
- The uts namespace: Isolating kernel and version identifiers. (UTS: Unix Timesharing System).

Docker Simplest Example


```
$ sudo docker run -i -t debian /bin/bash
```

```
$ docker run debian echo hello-world  
hello-world
```

- <http://containertutorials.com/>

Docker Notes

- Docker Object
 - Images
 - `docker build`
 - Containers
 - `docker run -i -t ubuntu /bin/bash`
- Will pull from configured registries
- Repositories and Registries
 - Git analogy (repo vs github)
 - The thing to remember here is a Docker repository is a place for you to publish and access your Docker images. Just like GitHub is a place for you to publish and access your git repos
 - Push/pull
- Need to do everything as root



Understanding Images

- Registry vs local Repository
 - Public (dockerhub) vs Private registries
 - docker search alpine-apache
- docker images
 - Image variants [image]:[tag]
 - docker image
- docker build
- docker pull

Pull the alpine image,

```
$ docker pull alpine
```

Check IP Address of the container

```
$ docker run alpine ifconfig
```

Launching a bash shell

```
$ docker run -i -t alpine /bin/bash
```

Docker Build Example

1. Create a Dockerfile

```
FROM smebberson/alpine-apache
ADD ./public-html/myindex.html /var/www/localhost/htdocs
```

2. Create a directory public_html with the following content in myindex.

```
<html>
<body>
Hi There - Static page served by Apache Server
</body>
</html>
```

3. Your directory should look like this

```
$ tree .
.
├── Dockerfile
└── public-html
    └── myindex.html
```

4. Create a Docker image

```
$ docker build -t my-apache2-alpine .
```

This will create a my-apache2 image.

5. Create a Docker Container running this image

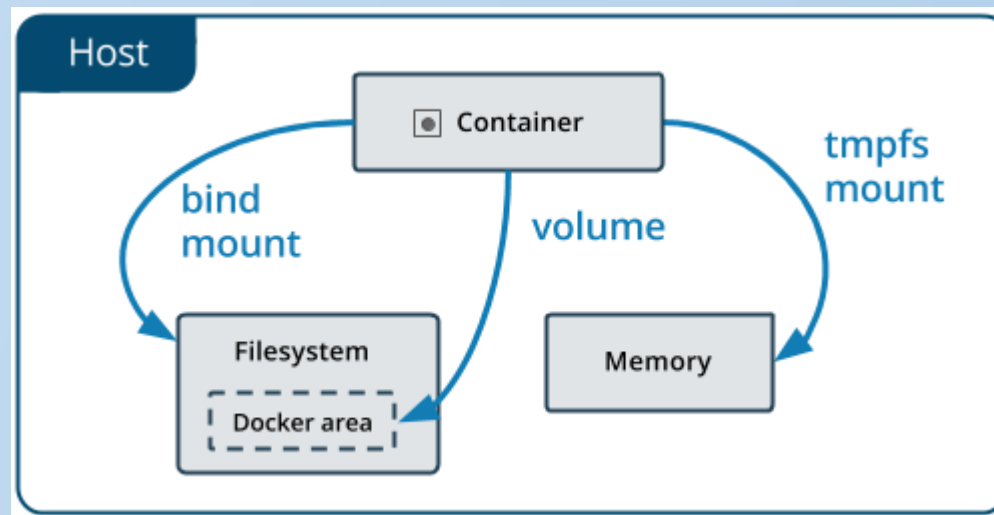
```
docker run -p 80:80 --name my-apache2-alpine-1 my-apache2-alpine
```

Dockerfile

- Commands
 - FROM
 - ADD
 - RUN
 - COPY
 - EXPOSE
 - ENVIRONMENT

Volumes and Network Port Binding

- Mount Host Directory as a 'volume'
 - Ugh overloaded term.. hear 'volume' means a rooted directory tree.

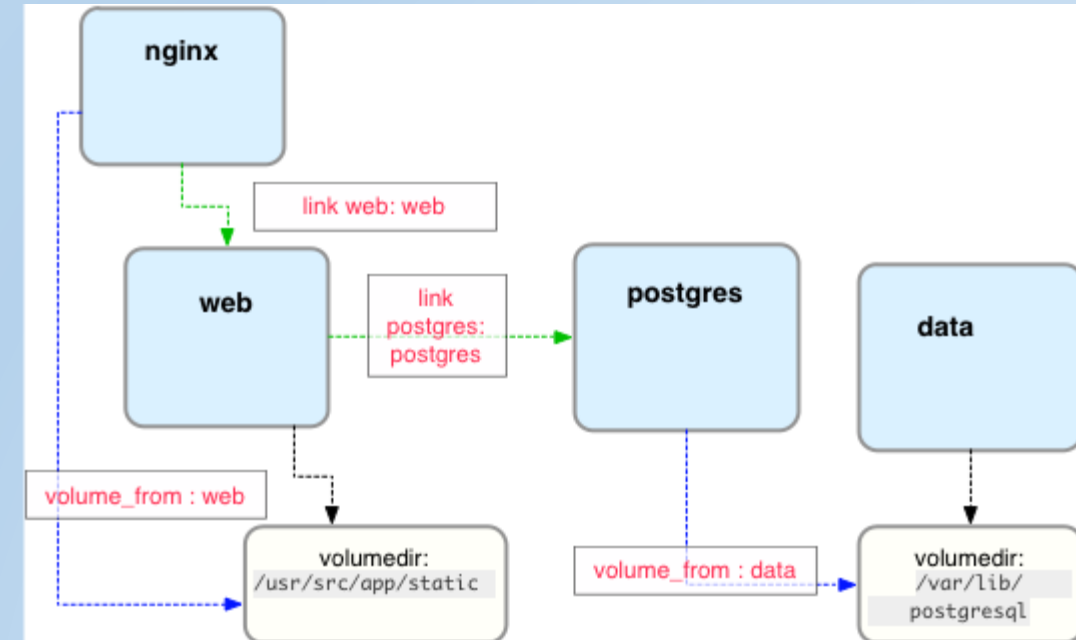


- Networking
 - Bind Host port to container port

Docker Compose

- http://containertutorials.com/linked/docker_compose.html

- Remember Flask servers are simple
 - Examples
 - Docker Compose with 2 ports/base



- <http://containertutorials.com/docker-compose/spring-boot-app.html>

Examples

- From

1. Create a Dockerfile with the following content

```
FROM debian:wheezy

RUN apt-get update && apt-get install -y cowsay fortune
```

2. Go to the directory containing Dockerfile and execute the following command to build an image

```
$ docker build -t test/cowsay-dockerfile .
```

You will see output as shown below

```
Sending build context to Docker daemon 2.048 kB
Sending build context to Docker daemon
Step 0 : FROM debian:wheezy
wheezy: Pulling from debian
7a3e804ed6c0: Pull complete
b96d1548a24e: Already exists

Status: Downloaded newer image for debian:wheezy
---> b96d1548a24e
Step 1 : RUN apt-get update && apt-get install -y cowsay fortune
---> Running in 4404353a3643
Get:1 http://security.debian.org wheezy/updates Release.gpg [1554 B]
Get:2 http://security.debian.org wheezy/updates Release [102 kB]
Get:3 http://httpredir.debian.org wheezy Release.gpg [2390 B]
.....
Setting up perl (5.14.2-21+deb7u2) ...
update-alternatives: using /usr/bin/prename to provide /usr/bin/rename
---> ca3618d10f2a
Removing intermediate container 4404353a3643
Successfully built ca3618d10f2a
```

3. Check that image has been created

REPOSITORY	TAG	IMAGE ID	CREATED	VIRTUAL SIZE
test/cowsay-dockerfile	latest	ca3618d10f2a	3 minutes ago	126.9 MB
docker-dev	dry-run-test-2	db155754d7fc	6 days ago	1.571 GB
<none>	<none>	b01392d005bb	6 days ago	1.571 GB
debian	wheezy	b96d1548a24e	7 days ago	84.97 MB
debian	latest	df2a0347c9d0	7 days ago	125.2 MB
dockerswarm/dind-master	latest	bb4cd757411e	7 days ago	159 MB
<none>	<none>	f672d2db20f6	7 days ago	1.571 GB
<none>	<none>	1fe07c1fdf52	8 days ago	1.571 GB
dockerswarm/swarm-test-env	latest	01e6a0da0825	2 weeks ago	515.5 MB
ubuntu	14.04	07f8e8c5e660	3 weeks ago	188.3 MB
hello-world	latest	91c95931e552	5 weeks ago	910 B
busybox	latest	8c2e06607696	5 weeks ago	2.433 MB

4. Run the cowsay program using the built image

```
$ docker run test/cowsay-dockerfile /usr/games/cowsay "Hi!"
```

This will execute and show the output

```
< Hi! >
  -----
      \   ^__^
       (oo)\_____)
        (_____)  )\/\
           ||----w |
           ||     ||
```

5. Removing a Docker Image : Docker image can be removed using the following command

```
$ docker rmi test/cowsay-dockerfile
```

Docker Commands

- http://containertutorials.com/get_started
- `docker <cmd>`
 - `info`
 - `run`
 - `run -it`
 - `attach`
 - `ps`
 - `start`
 - `stop`
 - `logs`
 - `rm`
 - `pause/unpause`
 - `network ls`
 - `rename`

Unravelling the Kubernetes Storage Abstraction

- There are Many Confusing Kubernetes Objects
 - Pod
 - Deployment, StatefulSet, DaemonSet
 - Volume
 - PersistentVolumeClaim
 - PersistentVolume
 - PersistentVolumeAttachment
 - StorageClass
 - CSIDriver


Where are the Containers?

- Pod
 - One or more Containers Running with a shared 'localhost'
- Deployment
 - A Template for deploying several Pods over multiple nodes (interchangeably)
- StatefulSet
 - A Template for deploying several Pods over multiple nodes (sticky id)
- DaemonSet
 - Template for running several Pods one per Node (or subset of Nodes)

Persistent Storage Objects

- Volume
- PersistentVolumeClaim
- PersistentVolume
- PersistentVolumeAttachment
- StorageClass
- CSIDriver

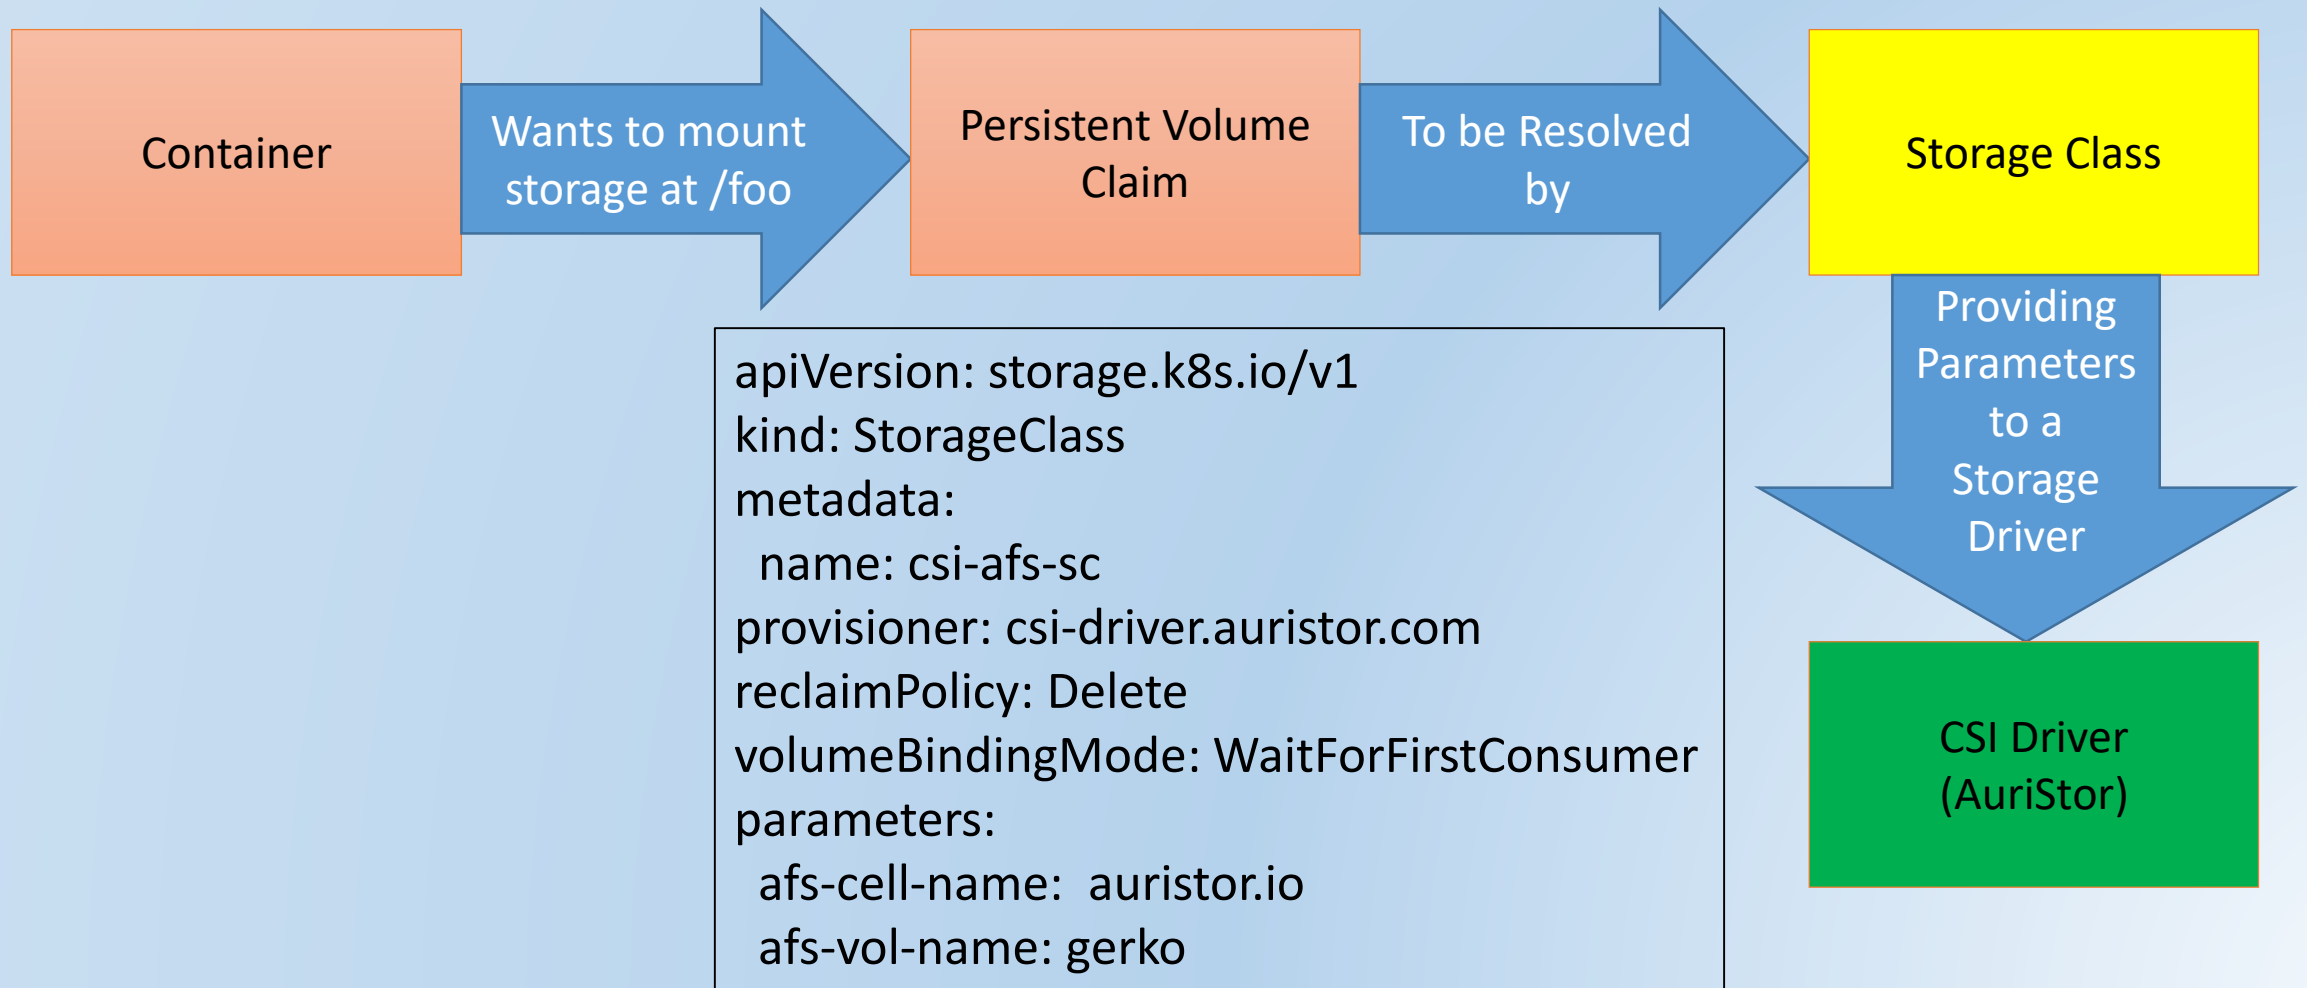
What the Application Developer Specifies



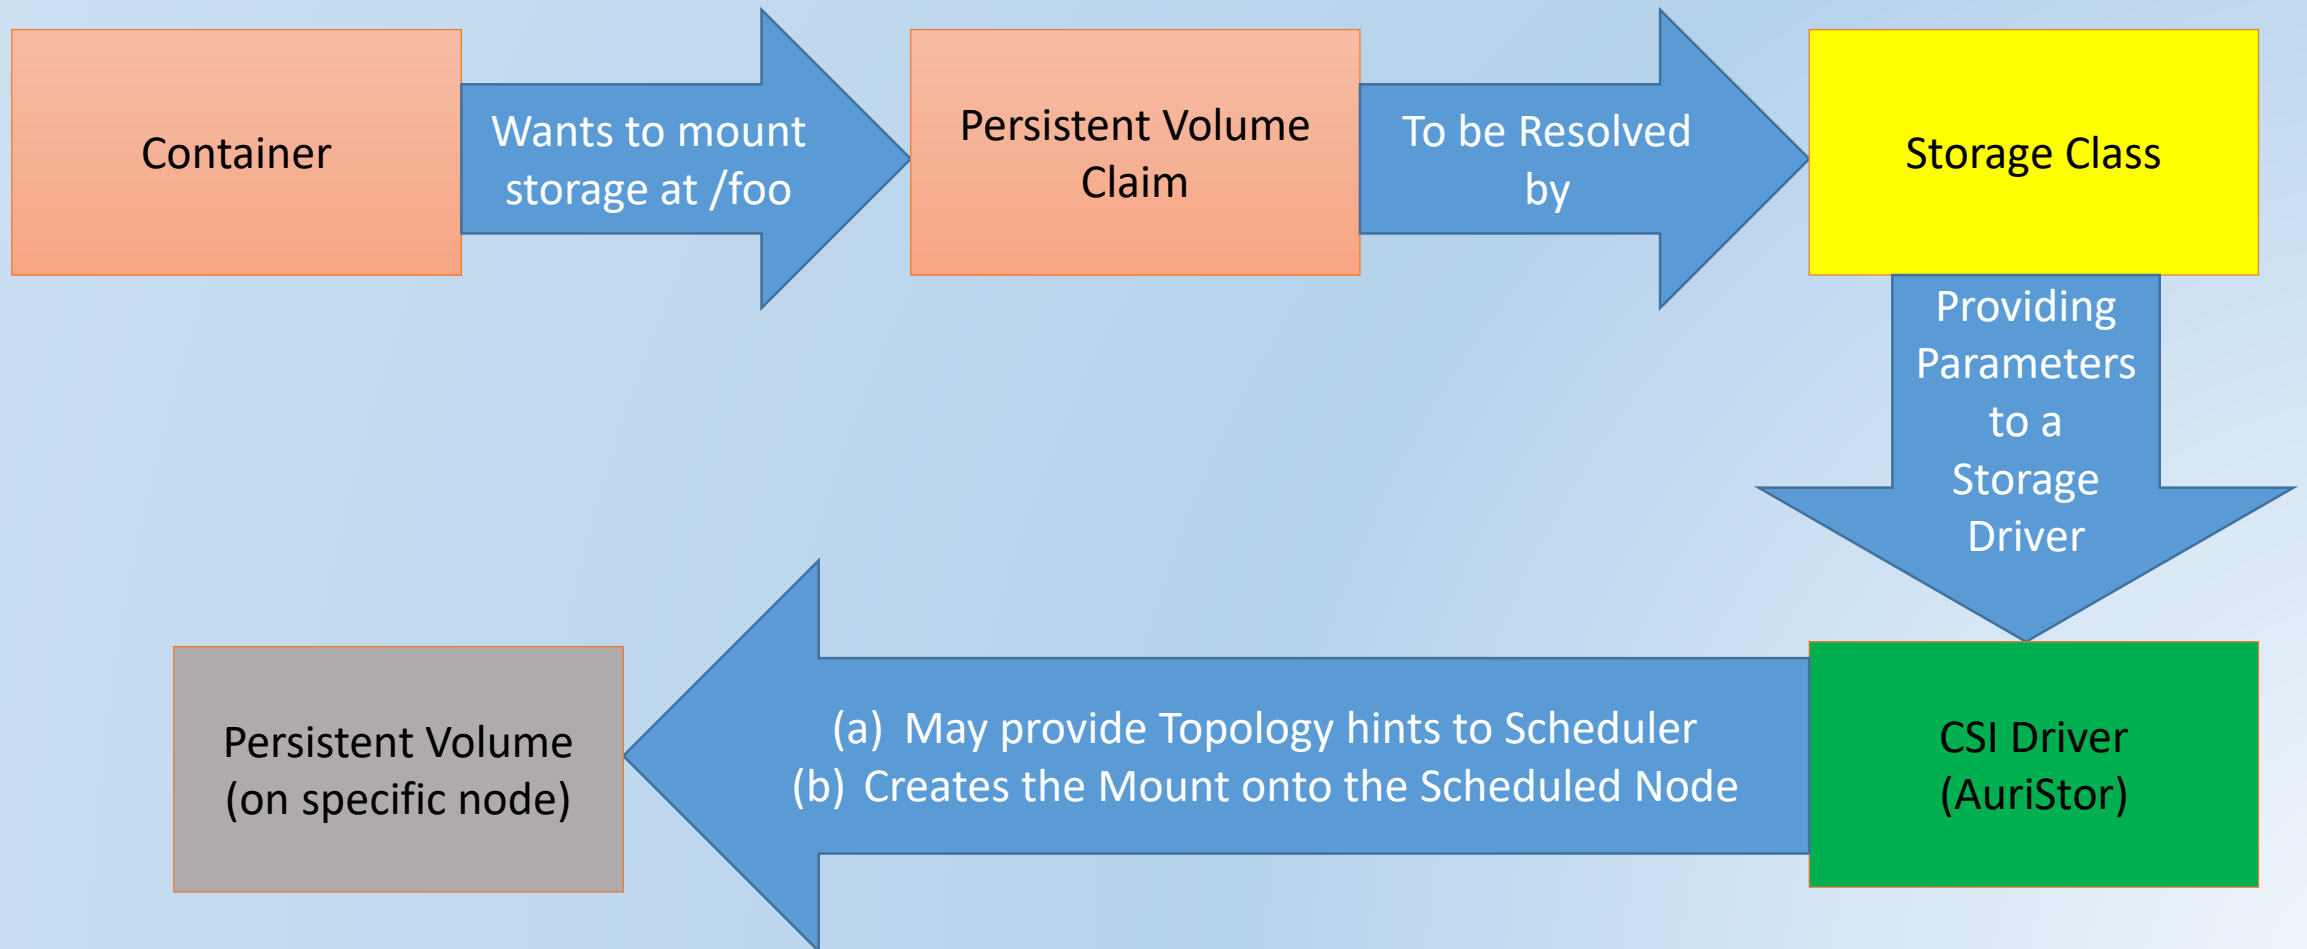
```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: csi-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  storageClassName: csi-afs-sc
```

```
kind: Pod
apiVersion: v1
metadata:
  name: my-csi-app
spec:
  containers:
    - name: my-frontend
      image: busybox
      volumeMounts:
        - mountPath: "/data"
          name: my-csi-volume
      command: [ "sleep", "1000000" ]
  volumes:
    - name: my-csi-volume
      persistentVolumeClaim:
        claimName: csi-pvc
```

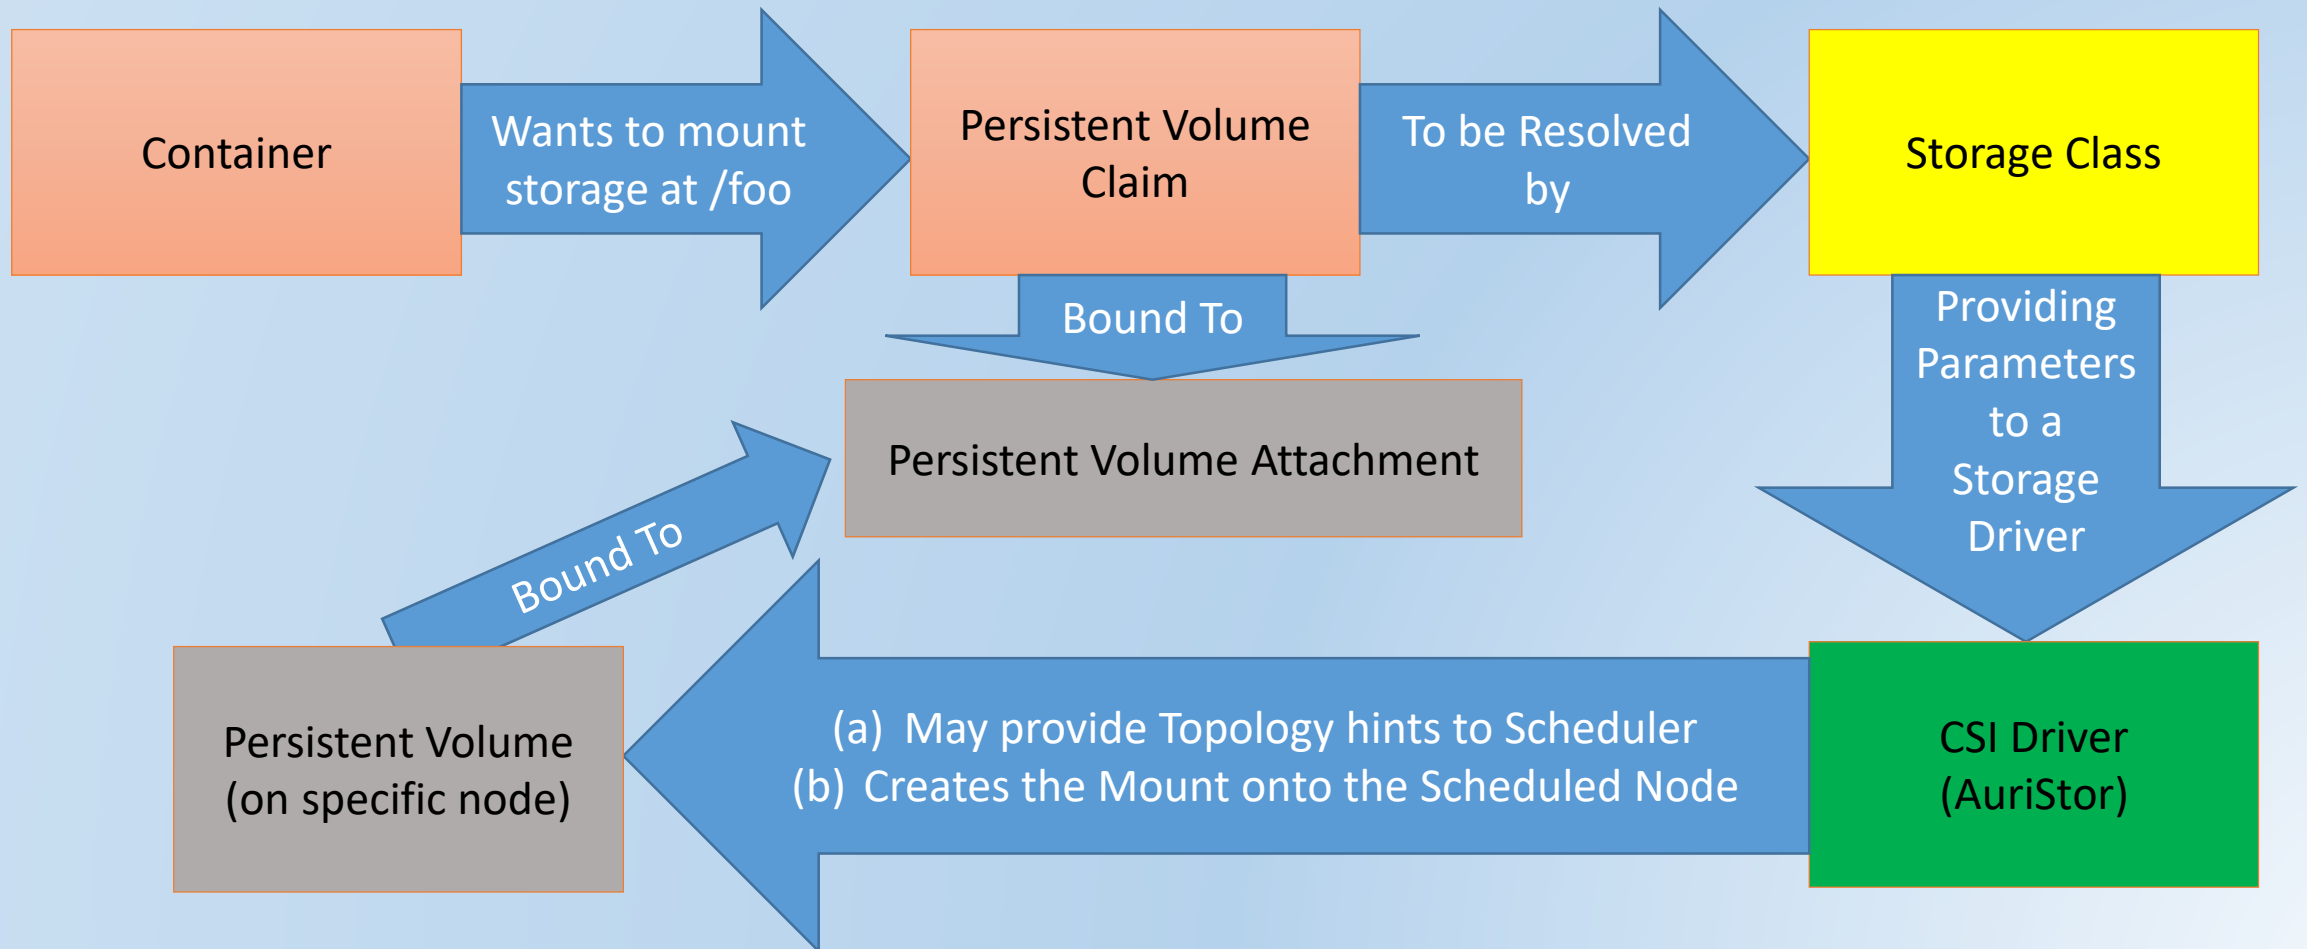
What the Cluster (Storage) admin Specifies



What the Kubernetes CSI Driver Does



What the external-attacher does



The Storage Class Parameters to CSI Driver

- AuriStor Driver Specific
 - Volume:
 - Existing Cell + Volume
 - Scratch + Quota
 - Topology
 - Don't care where the volume is
 - Schedule 'near' a File Server (ie Rack) with Volume (Replicate if necessary)
 - Do not Start until replication completes
 - Start immediately if none, but also start replication

Distributed Systems Study Group

First Meeting: Dec 3rd, 2019

Questions?

Gerry Seidman
gerry@iamx.com